COLUMBIA UNIVERSITY

Department of Economics

Completed in Partial Fulfillment of the Requirements of

ECON UN3412

Introduction to Econometrics

Spring 2026

Homework 5

Seyhan Erden

seyhan.erden@columbia.edu

Tyler Sotomayor

tyler.sotomayor@columbia.edu

Special thanks to the Teaching Assistants of ECON UN3412

Daniel Garcia

dg3342@columbia.edu

Chris Kye

mk4449@columbia.edu

Shivesh Babu Narayan

sb5211@columbia.edu

Luca Salvatori

ls4241@columbia.edu

Zijian Zhang

zz3263@columbia.edu

Contents

Problem 1: Concealed Weapons Laws and Violent Crime 1

Problem 2: Panel Regression Tables (ln(vio), ln(rob), ln(mur)) 9

Problem 3: Stock Market Performance and Firm Returns 18

Ungraded Exercises 30

Appendix: Tables and Figures 45

Replication Files 48

ECON UN3412 Spring 2026: Homework 5 Tyler Sotomayor

Introduction to Econometrics tyler.sotomayor@columbia.edu

Instructor: Seyhan Erden Due Date: 8 April 2026

Problem 1 [30 points]

Some U.S. states have enacted laws that allow citizens to carry concealed weapons. These laws are known as "shall-issue" laws because they instruct local authorities to issue a concealed weapons permit to all applicants who are citizens, mentally competent, and have not been convicted of a felony (some states have some additional restrictions). Proponents argue that, if more people carry concealed weapons, crime will decline because criminals are deterred. Opponents argue that crime will increase because of accidental or spontaneous use of the weapon. In this exercise, you will analyze the effect of concealed weapons laws on three different categories of crimes: violent crimes; robberies (such as the robbery of a convenience store); and murder (many of which are spontaneous acts of passion).

Data file and the description of variables are at the end of this problem set. Estimate

(i) [3 points] A regression of $\ln(vio)$ against *shall*.

SOLUTION:

The baseline specification regresses the natural log of the violent crime rate on the shall-issue indicator alone, using heteroskedasticity-robust (Eicker–Huber–White) standard errors. In Stata:

```
use "handguns.dta", clear
```

```
gen ln_vio = ln(vio)
reg ln_vio shall, robust
```

The estimated equation is:

$$\boxed{\widehat{\ln(vio)}_i = \underset{(0.019)}{6.135} - \underset{(0.048)}{0.443}\, shall_i}$$

$n = 1{,}173, \quad R^2 = 0.0866.$

The coefficient on *shall* is -0.443 and is statistically significant at the 1% level ($t = -9.32$, $p < 0.001$). Without any controls, states with shall-issue laws in effect have, on average, approximately 44.3% lower violent crime rates than states without such laws. However, this naive estimate likely suffers from omitted variable bias (OVB), as states that adopt shall-issue laws may differ systematically from non-adopting states in ways that also affect crime. See Table 4 in the Appendix for the full regression output. ■

(ii) [3 points] A regression of $\ln(vio)$ against *shall*, *incarc_rate*, *density*, *avginc*, *pop*, *pb1064*, *pw1064*, and *pm1029*.

SOLUTION:

The multiple regression adds seven state-characteristic control variables that capture cross-state differences in incarceration policy, urbanization, income, population size, and demographic composition. In Stata:

```
reg ln_vio shall incarc_rate density avginc pop ///
    pb1064 pw1064 pm1029, robust
```

The estimated equation is:

$$\begin{aligned}\widehat{\ln(vio)}_i = \underset{(0.609)}{2.982} &- \underset{(0.035)}{0.368}\, shall_i + \underset{(0.0002)}{0.0016}\, incarc_rate_i \\ &+ \underset{(0.014)}{0.027}\, density_i + \underset{(0.007)}{0.001}\, avginc_i + \underset{(0.003)}{0.043}\, pop_i \\ &+ \underset{(0.020)}{0.081}\, pb1064_i + \underset{(0.010)}{0.031}\, pw1064_i \\ &+ \underset{(0.012)}{0.009}\, pm1029_i\end{aligned}$$

$n = 1{,}173, \quad R^2 = 0.5643, \quad \bar{R}^2 = 0.5613.$

The coefficient on *shall* is -0.368 and remains highly significant ($t = -10.59$, $p < 0.001$). The inclusion of controls substantially raises R^2 from 0.087 to 0.564, indicating that the control variables explain a large share of the cross-state variation in violent crime. See Table 5 in the Appendix for the full regression output. ■

(a) [8 points] Interpret the coefficient on *shall* in regression (ii). Is this estimate large or small in a "real-world" sense?

SOLUTION:

The coefficient on *shall* in regression (ii) is $\hat{\beta}_{shall} = -0.368$. Because the dependent variable is $\ln(vio)$ and *shall* is a binary indicator, this coefficient has a log-level (semi-elasticity) interpretation: holding all other regressors fixed, the presence of a shall-issue law is associated with a change of -0.368 in the log of the violent crime rate. By the standard log-point approximation, this corresponds to approximately a $100 \times (-0.368) = -36.8\%$ reduction in the violent crime rate

(Stock and Watson 2020, Ch. 8). More precisely, using the exact transformation $\exp\left(\hat{\beta}\right) - 1 = \exp(-0.368) - 1 \approx -0.308$, the estimated reduction is about 30.8%.

Real-world magnitude. Given that the sample mean of the violent crime rate is approximately 503 incidents per 100,000 population, a 30–37% reduction would imply roughly 155–185 fewer violent crimes per 100,000 residents. This is an extraordinarily large effect. To put it in context, the entire decline in the U.S. violent crime rate from its 1991 peak to 1999 was on the order of 30%, and this decline is attributed to a wide range of factors (increased incarceration, demographic shifts, economic growth, policing strategies, among others). It would be implausible that shall-issue laws alone could account for an effect of this magnitude. This strongly suggests that the pooled OLS estimate, even after controlling for observables, still suffers from substantial omitted variable bias — likely because states that adopt shall-issue laws differ from non-adopting states in unobserved ways (e.g., political culture, historical crime trends, law enforcement practices) that are correlated with both the passage of the law and with crime rates. The original and influential study by Lott and Mustard (1997) reported large crime-reducing effects of shall-issue laws using similar cross-state variation, but these estimates have been widely critiqued for sensitivity to specification choices (Donohue and Ayres, 2003; National Research Council, 2005).

$$\boxed{\hat{\beta}_{shall} = -0.368 \implies \approx 30.8\%\text{–}36.8\% \text{ reduction; implausibly large.}}$$

■

Reflections: The log-level interpretation of a dummy variable coefficient follows from the general result that for $\ln Y = \beta_0 + \beta_1 D + \mathbf{X}'\boldsymbol{\gamma} + u$, the exact percentage change associated with D switching from 0 to 1 is $100 \times [\exp\left(\hat{\beta}_1\right) - 1]\%$ (see Wooldridge 2019, §6.2; Stock and Watson 2020, §8.1). The $100 \times \hat{\beta}_1$ approximation works well for small $|\hat{\beta}_1|$ but becomes less accurate as $|\hat{\beta}_1|$ grows. For $\hat{\beta}_1 = -0.368$, the discrepancy between the approximation (-36.8%) and the exact figure (-30.8%) is notable, illustrating why the exact calculation is preferred when the coefficient is large.

The implausibly large magnitude of the OLS estimate motivates the use of panel data methods in Problem 2, where state and year fixed effects will be introduced to control for unobserved heterogeneity. This is the central methodological lesson of this exercise: cross-sectional OLS can be severely biased when unobserved confounders are present, and fixed effects estimation provides a way to address time-invariant omitted variables.

(b) [8 points] Does adding the control variables in regression (ii) change your conclusions about the effect of a shall-carry law, relative to regression (i)?

Solution:

Adding the seven state-characteristic control variables changes the magnitude of the estimated effect but does not change the qualitative conclusion. In regression (i), the coefficient on *shall* is -0.443 (SE $= 0.048$), while in regression (ii) it becomes -0.368 (SE $= 0.035$). Both estimates are negative and highly statistically significant at conventional levels.

The coefficient shrinks by about 17% in absolute value, from -0.443 to -0.368,

when the controls are added. This attenuation is consistent with positive omitted variable bias in regression (i): states without shall-issue laws tend to have characteristics associated with higher crime (e.g., higher population density, larger black population share, higher incarceration rates), and controlling for these factors removes part of the spurious negative correlation between *shall* and $\ln(vio)$.

However, the estimated effect remains very large—a 30–37% reduction in violent crime is still implausibly attributable to a single policy change. The fact that the estimate barely changes (from -0.443 to -0.368) despite a large increase in R^2 (from 0.087 to 0.564) suggests that the included controls, while explaining much of the variation in crime, do not adequately address the confounding that generates the large negative *shall* coefficient. There are likely important omitted variables that are correlated with both shall-issue adoption and crime rates, which the observable controls cannot capture. This motivates the use of state fixed effects in Problem 2 to control for all time-invariant state-specific factors.

$$\boxed{\text{Qualitative conclusion unchanged; magnitude decreases modestly: } -0.443 \to -0.368.}$$

■

(c) [8 points] Suggest a variable which varies across states but plausibly varies little, or not at all, over time, and which plausibly could cause omitted variable bias in regression (ii).

SOLUTION:

A strong candidate is the **prevailing gun culture or attitudes toward firearms ownership** in each state. This variable reflects deeply rooted his-

torical, cultural, and geographic factors—such as whether a state has a frontier heritage, rural traditions of hunting, or a longstanding culture of individual gun ownership—that vary substantially across states but change very slowly (if at all) over a 23-year panel.

Gun culture plausibly satisfies both conditions required for omitted variable bias (Stock and Watson 2020, Ch. 6): it must be correlated with the included regressor (*shall*) and with the dependent variable ($\ln(vio)$), conditional on the other controls.

First, states with a stronger gun culture are more likely to pass shall-issue laws, because both the political demand for such legislation and the legislative willingness to enact it are higher where gun ownership is culturally normalized. This implies $\text{Corr}(gun\ culture, shall) > 0$.

Second, gun culture is plausibly correlated with violent crime rates through multiple channels. States with deep gun cultures may have higher rates of gun ownership, which could increase or decrease crime depending on one's theoretical priors. More importantly, such states also tend to differ in other hard-to-measure ways (e.g., social norms about conflict resolution, rural vs. urban composition beyond what *density* captures, quality of law enforcement). These factors may independently affect crime rates.

Because gun culture is positively correlated with *shall* and plausibly negatively correlated with crime (rural, gun-culture states tend to have lower violent crime rates for reasons beyond just gun laws), omitting it would bias $\hat{\beta}_{shall}$ downward (more negative), making shall-issue laws appear more effective at reducing crime than they truly are. This is a form of negative omitted variable bias (since

$\text{Corr}(\textit{gun culture}, \textit{shall}) > 0$ and the omitted variable is negatively correlated with the error term from the perspective of inflating the negative coefficient). This variable is approximately time-invariant within the sample period, which is precisely the type of confounding that state fixed effects are designed to eliminate.

Gun culture / attitudes toward firearms (varies across states, nearly time-invariant).

■

REFLECTIONS: The omitted variable bias formula for the simple case is

$$\text{Bias}(\hat{\beta}_1) = \beta_2 \cdot \frac{\text{Cov}(X_1, X_2)}{\text{Var}(X_1)},$$

where X_2 is the omitted variable and β_2 is its true coefficient (Wooldridge 2019, §3.3; Stock and Watson 2020, §6.1). This exercise is designed to motivate panel data fixed effects: variables that are roughly constant within each state over time are exactly the type of confounders that entity (state) fixed effects absorb. The "within" estimator differences out all time-invariant state characteristics, whether observed or unobserved, eliminating OVB from such sources.

Problem 2 [30 points]

Read the notes to Tables 1–3, then fill in the tables.

Table 1

The Effect of Concealed Handgun Laws on Violent Crime: Regression Results

Dependent variable: $\ln(vio)$

	(1)	(2)	(3)	(4)	(5)
Coefficient on *shall*					
	( )	( )	( )	( )	( )
State characteristic control variables[a]?	no	yes	yes	yes	yes
State fixed effects?	no	no	yes	yes	yes
Year fixed effects?	no	no	no	yes	yes
F-statistic: state FE = 0	–	–	( )	( )	–
F-statistic: year FE = 0	–	–	–	( )	( )
HAC (clustered) SEs?	no	no	no	no	yes
n					

Table 2

The Effect of Concealed Handgun Laws on Robberies: Regression Results

Dependent variable: $\ln(rob)$

Same structure as Table 1. See the notes to Table 1.

Table 3

The Effect of Concealed Handgun Laws on Murders: Regression Results

Dependent variable: $\ln(mur)$

Same structure as Table 1. See the notes to Table 1. $n = 1173$.

Notes: All regressions include an intercept. For regressions (1)–(4), heteroskedasticity-robust standard errors appear in parentheses below estimated coefficients; for regression (5), the standard errors are heteroskedasticity-robust and clustered at the state level, so as to allow for serial correlation in the error within a state. p-values appear in parentheses beneath heteroskedasticity-robust F-statistics (or, for regression (5), heteroskedasticity-robust-clustered F-statistic).

[a]Regressions with "state characteristic control variables" include the following regressors: *incarc_rate*, *density*, *avginc*, *pop*, *pb1064*, *pw1064*, *pm1029*.

SOLUTION:

The five specifications estimated for each dependent variable are:

(1) Pooled OLS of $\ln(Y)$ on *shall* only, with heteroskedasticity-robust standard errors.

(2) Pooled OLS adding the seven state-characteristic control variables (*incarc_rate*, *density*, *avginc*, *pop*, *pb1064*, *pw1064*, *pm1029*), with robust SEs.

(3) Adding 50 state fixed effect dummies (with one state as the omitted reference category), plus controls, with robust SEs. State FE absorb all time-invariant unobserved state-level heterogeneity. Identification of $\hat{\beta}_{shall}$ comes exclusively from *within-state* variation: states that changed their shall-issue status during 1977–1999.

(4) Adding 22 year fixed effect dummies (with one year as reference) to specification (3), with robust SEs. Year FE absorb aggregate time trends common to all states (e.g., national crime trends, federal policy changes, macroeconomic cycles). This is the canonical two-way fixed effects (TWFE) model (Stock and Watson 2020, Ch. 10).

(5) Same as (4), but with standard errors clustered at the state level via `vce(cluster stateid)`. Clustering allows for arbitrary serial correlation (autocorrelation) in the error term within each state over time, as well as heteroskedasticity, producing HAC-robust inference (Arellano 1987; Stock and Watson 2020, §10.5).

The Stata code for the full set of regressions (taking *ln_vio* as an example) proceeds as follows:

```
* Generate log dependent variables
gen ln_vio = ln(vio)

* Create state and year dummies
tabulate stateid, generate(st)
tabulate year, generate(yr)

* Column (1): shall only, robust SEs
reg ln_vio shall, robust

* Column (2): + controls, robust SEs
reg ln_vio shall incarc_rate density avginc pop ///
    pb1064 pw1064 pm1029, robust

* Column (3): + state FE, robust SEs
reg ln_vio shall incarc_rate density   ///
    avginc pop pb1064 pw1064 pm1029    ///
```

```
        st2-st51, robust
    testparm st2-st51  // F-test: state FE

    * Column (4): + state FE + year FE
    reg ln_vio shall incarc_rate density   ///
        avginc pop pb1064 pw1064 pm1029    ///
        st2-st51 yr2-yr23, robust
    testparm st2-st51  // F-test: state FE
    testparm yr2-yr23  // F-test: year FE

    * Column (5): + clustered SEs
    reg ln_vio shall incarc_rate density   ///
        avginc pop pb1064 pw1064 pm1029    ///
        st2-st51 yr2-yr23,                 ///
        vce(cluster stateid)
    testparm yr2-yr23  // F-test: year FE
```

The completed tables follow below. The F-statistics test the null hypothesis that all fixed effect coefficients are jointly zero; the p-values appear in parentheses beneath.

Table 1: The Effect of Concealed Handgun Laws on Violent Crime: Regression Results

Dependent variable: $\ln(vio)$

	(1)	(2)	(3)	(4)	(5)
Coefficient on *shall*	-0.443***	-0.368***	-0.046**	-0.028	-0.028
	(0.048)	(0.035)	(0.020)	(0.019)	(0.042)
State characteristic control variables[a]?	no	yes	yes	yes	yes
State fixed effects?	no	no	yes	yes	yes
Year fixed effects?	no	no	no	yes	yes
F-statistic: state FE = 0	–	–	210.38	309.29	–
			(0.000)	(0.000)	
F-statistic: year FE = 0	–	–	–	13.90	20.67
				(0.000)	(0.000)
HAC (clustered) SEs?	no	no	no	no	yes
n	1173	1173	1173	1173	1173

Notes: All regressions include an intercept. For regressions (1)–(4), heteroskedasticity-robust standard errors appear in parentheses below estimated coefficients; for regression (5), the standard errors are heteroskedasticity-robust and clustered at the state level, so as to allow for serial correlation in the error within a state. p-values appear in parentheses beneath heteroskedasticity-robust F-statistics (or, for regression (5), heteroskedasticity-robust-clustered F-statistic). $^{***}p < 0.01$; $^{**}p < 0.05$; $^{*}p < 0.10$.
[a]Regressions with "state characteristic control variables" include the following regressors: *incarc_rate*, *density*, *avginc*, *pop*, *pb1064*, *pw1064*, *pm1029*.

Table 2: The Effect of Concealed Handgun Laws on Robberies: Regression Results

Dependent variable: $\ln(rob)$

	(1)	(2)	(3)	(4)	(5)
Coefficient on *shall*	-0.773***	-0.529***	-0.008	0.027	0.027
	(0.069)	(0.051)	(0.026)	(0.025)	(0.053)
State characteristic control variables[a]?	no	yes	yes	yes	yes
State fixed effects?	no	no	yes	yes	yes
Year fixed effects?	no	no	no	yes	yes
F-statistic: state FE = 0	–	–	190.47	243.39	–
			(0.000)	(0.000)	
F-statistic: year FE = 0	–	–	–	12.39	24.73
				(0.000)	(0.000)
HAC (clustered) SEs?	no	no	no	no	yes
n	1173	1173	1173	1173	1173

Notes: See the notes to Table 1.

Table 3: The Effect of Concealed Handgun Laws on Murders: Regression Results

Dependent variable: ln(*mur*)

	(1)	(2)	(3)	(4)	(5)
Coefficient on *shall*	-0.473***	-0.313***	-0.061**	-0.015	-0.015
	(0.049)	(0.036)	(0.027)	(0.027)	(0.039)
State characteristic control variables[a]?	no	yes	yes	yes	yes
State fixed effects?	no	no	yes	yes	yes
Year fixed effects?	no	no	no	yes	yes
F-statistic: state FE = 0	–	–	88.22	106.69	–
			(0.000)	(0.000)	
F-statistic: year FE = 0	–	–	–	9.73	18.75
				(0.000)	(0.000)
HAC (clustered) SEs?	no	no	no	no	yes
n	1173	1173	1173	1173	1173

Notes: See the notes to Table 1.

The key takeaway across all three tables is striking and instructive. The pooled OLS estimates in columns (1) and (2) suggest that shall-issue laws are associated with large, statistically significant reductions in violent crime, robbery, and murder. However, once state fixed effects are introduced in column (3), the estimated effects shrink dramatically—by roughly 87–99%—indicating that the cross-sectional estimates were severely contaminated by omitted time-invariant state characteristics. Adding year fixed effects in column (4) further reduces the coefficients, and for all three crime categories the estimates become statistically insignificant. In column (5), where standard errors are clustered at the state level to allow for serial correlation within states, the

estimates remain insignificant and the standard errors approximately double compared to column (4). This illustrates the importance of accounting for serial correlation in panel data, as emphasized by Bertrand et al. (2004). The final conclusion from the preferred specification (column 5) is that there is no statistically significant evidence that shall-issue laws affect violent crime, robbery, or murder rates.

After state and year FE with clustered SEs, shall-issue laws have no statistically significant effect on crime.

■

Reflections: The dramatic change in the *shall* coefficient across specifications illustrates the core logic of fixed effects estimation. The pooled OLS estimator exploits both *between-state* and *within-state* variation, but the between-state variation is contaminated by unobserved state-level confounders. The within (fixed effects) estimator uses only within-state variation over time, effectively comparing each state to itself before and after the law change, thereby eliminating all time-invariant confounding.

The F-statistics for state fixed effects are extremely large (e.g., 210–309 for violent crime), confirming that states differ enormously in their baseline crime rates—differences that the pooled regressions attribute partly to *shall*. The year fixed effects are also highly significant, reflecting common national crime trends (the crime wave of the late 1980s and the subsequent decline). By including both sets of fixed effects, the TWFE model identifies the *shall* effect from the *timing* of law adoption relative to other states, after removing both state-level means and common time shocks.

The increase in standard errors from column (4) to column (5) demonstrates why clustering is essential in panel data. When errors are serially correlated within states,

conventional robust SEs understate the true sampling uncertainty, leading to spuriously small p-values. Clustering at the entity level (Arellano 1987; Bertrand et al. 2004; Stock and Watson 2020, §10.5) corrects this by allowing arbitrary within-cluster correlation.

Problem 3 [40 points]

In order to determine the effects of stock market performance on individual firm's performance, you collect a daily frequency data on firm's performance for n sample of publicly listed companies for the period of January 1, 2021 – January 1, 2026 on a (trading) daily frequency.

(a) [8 points] What measures of firm performance would you include in an equation? What are some of the timing issues?

SOLUTION:

The most natural measure of firm performance in the context of stock market analysis is the **daily log return**, defined as

$$r_{it} = 100 \times \big[\ln(P_{it}) - \ln(P_{i,t-1})\big],$$

where P_{it} is the adjusted closing price of firm i on trading day t. The log-difference formulation is the standard econometric approach to measuring percentage changes (Stock and Watson 2020, Ch. 14): applying a log transformation to the price level and then first-differencing yields a variable that is (i) approximately equal to the simple percentage return $(P_{it} - P_{i,t-1})/P_{i,t-1}$ for small daily changes, (ii) additive over time (multi-day log returns are the sum of daily log returns), and (iii) approximately stationary—an essential property for valid regression inference, since levels of stock prices are typically non-stationary (exhibiting unit roots). The adjusted close price is preferred over the raw close because it accounts for stock splits, dividends, and other corporate actions that mechanically alter the price level without reflecting true performance changes.

Alternative measures one might consider include: (i) **abnormal returns** (the firm's return minus the expected return from an asset pricing model), which isolate the firm-specific component of performance; (ii) **simple percentage returns**, $100 \times (P_{it} - P_{i,t-1})/P_{i,t-1}$, which are nearly identical to log returns for daily data but lack the additive property; and (iii) **trading volume** or **market capitalization changes**, though these capture liquidity and size rather than performance per se.

Several timing issues arise with daily financial data. First, there is the question of **non-synchronous trading**: if some stocks close trading at slightly different times (e.g., due to trading halts or different exchange closing times), daily returns may not reflect the same information set, introducing measurement error. Second, stock prices may react to market-wide news with a **lag**—particularly for less liquid, smaller-cap stocks—which would argue for including lagged market returns alongside the contemporaneous market return in the regression (the Dimson 1979 correction). Third, prices respond to information released **between trading days** (overnight and weekend effects), so the daily close-to-close return conflates intraday movements and overnight reactions. Fourth, announcements or events (earnings releases, Federal Reserve decisions) typically occur at specific times; if the event falls after market close, the price impact appears on the following trading day, creating a one-day lead-lag alignment issue. Finally, the choice of whether to use **calendar time** or **trading-day time** matters: regressions should be run on trading days only, as weekends and holidays introduce spurious gaps.

$$\boxed{r_{it} = 100 \times \left[\ln(P_{it}) - \ln(P_{i,t-1})\right] \quad \text{(daily log return using adjusted close)}}$$

■

(b) [8 points] What other factors might you control for in the equation?

SOLUTION:

Beyond the overall stock market return (e.g., the S&P 500 return), several additional factors may influence individual firm stock returns and should be considered as controls, drawing on the asset pricing literature and the econometrics of panel data.

At the **firm level**, one might control for: (i) **firm size** (measured by log market capitalization), since small firms tend to exhibit higher average returns and greater sensitivity to market shocks than large firms—this is the "size effect" documented by Fama and French (1993); (ii) **book-to-market ratio**, capturing the "value premium" whereby high book-to-market ("value") firms earn higher average returns than low book-to-market ("growth") firms; (iii) **momentum** (past cumulative returns over the prior 3–12 months), since stocks that have performed well recently tend to continue outperforming in the short run (Jegadeesh and Titman 1993); and (iv) **trading volume or liquidity measures**, as illiquid stocks may exhibit larger price swings and lagged responses to market movements.

At the **macroeconomic and market level**, useful controls include: (v) **interest rate changes** (e.g., changes in the federal funds rate or the 10-year Treasury

yield), since monetary policy affects discount rates and hence stock valuations differently across sectors; (vi) **credit spreads** (the difference between corporate bond yields and risk-free rates), which capture changes in perceived default risk and economic uncertainty; (vii) **exchange rate fluctuations**, particularly relevant for firms with international operations; and (viii) **sector-specific indices or returns**, since firms within the same industry are exposed to common sector-level shocks (e.g., oil prices for energy firms, interest rates for banks) that are distinct from market-wide movements.

At the **time-series level**, controlling for (ix) **day-of-week effects** (the "Monday effect") and (x) **lagged dependent variables** (lagged firm returns) can account for well-documented calendar anomalies and return persistence or mean reversion. In a panel context, **firm fixed effects** (α_i) absorb all time-invariant firm characteristics (sector, size category, geographic location, management quality) and **time fixed effects** (δ_t) absorb all market-wide daily shocks (news, sentiment, macroeconomic releases), making them powerful catch-all controls.

Key controls: firm FE, time FE, size, book-to-market, momentum, interest rates, sector returns.

■

(c) [8 points] Write an equation that allows you to estimate the effects of the overall stock market performance on the percentage change in firm's stock price. How would you estimate this equation? Why would you choose this method?

SOLUTION:

The baseline equation for estimating the effect of overall market performance

on the percentage change in a firm's stock price in a panel setting is:

$$\boxed{\Delta \ln(P_{it}) = \alpha_i + \beta \cdot \Delta \ln(P_{mt}) + u_{it}, \quad i = 1, \ldots, n, \;\; t = 1, \ldots, T}$$

where $\Delta \ln(P_{it}) = \ln(P_{it}) - \ln(P_{i,t-1})$ is the daily log return of firm i on trading day t (the log-difference measure of percentage change), $\Delta \ln(P_{mt})$ is the daily log return of the S&P 500 index on day t, α_i is a firm-specific fixed effect (intercept), and u_{it} is the idiosyncratic error term. The parameter of interest is β, which measures the average sensitivity of firm stock returns to market-wide movements—the "market beta" in the Capital Asset Pricing Model (CAPM) tradition (Sharpe 1964; Lintner 1965).

This equation would be estimated using **fixed effects (within) estimation**. In Stata, this is implemented via the `xtreg` command with the `fe` option:

```
xtset firm_id trading_date
xtreg log_ret_firm log_ret_market, ///
    fe vce(robust)
```

The fixed effects estimator is chosen for several reasons. First, the panel structure of the data (multiple firms observed over the same trading days) naturally calls for a panel estimator rather than pooled OLS. Second, different firms have different average returns (reflecting differences in risk, industry, growth prospects, etc.), and these firm-specific mean returns are captured by α_i. If α_i is correlated with r_{mt}—which is plausible if, for example, firms with higher average returns are also more sensitive to market conditions—then pooled OLS would yield biased estimates of β due to omitted variable bias (Wooldridge 2019, Ch. 14). The

fixed effects estimator eliminates α_i via the within transformation (demeaning each variable by its firm-specific mean), producing consistent estimates of β even when $\text{Cov}(\alpha_i, r_{mt}) \neq 0$ (Stock and Watson 2020, Ch. 10).

Third, the within estimator is numerically equivalent to the Least Squares Dummy Variable (LSDV) approach—including a full set of n firm dummies—but is computationally more efficient when n is large. Fourth, using heteroskedasticity-robust standard errors (or clustered standard errors at the firm level) ensures valid inference even if the variance of u_{it} differs across firms or exhibits serial correlation within firms.

An extended specification that also controls for aggregate time effects would be:

$$\Delta \ln(P_{it}) = \alpha_i + \delta_t + \beta \cdot \Delta \ln(P_{mt}) + u_{it}$$

where δ_t are time (day) fixed effects. However, note that including day fixed effects alongside r_{mt} creates a near-perfect multicollinearity problem, since r_{mt} is constant across firms on any given day and is thus almost perfectly spanned by the day dummies. In practice, one would either include day FE *without* r_{mt} (letting the day dummies absorb all common daily shocks) or include r_{mt} without day FE. The choice depends on the research question: if the goal is to estimate β, one uses r_{mt} directly; if the goal is to control for all common daily shocks without parametric restriction, one uses day FE. ■

Reflections: This equation is essentially the market model from empirical finance, a single-factor version of the CAPM. In the CAPM framework, α_i should equal $r_f(1 - \beta_i)$ where r_f is the risk-free rate, and in an efficient market, the

"Jensen's alpha" (excess α_i) should be zero on average. The fixed effects estimator in a panel setting with a common regressor $\Delta \ln(P_{mt})$ is equivalent to running separate time-series regressions for each firm and then averaging the slope coefficients (this is the Fama and MacBeth 1973 intuition). The key econometric insight is that firm fixed effects control for all unobserved time-invariant firm characteristics, making the panel FE estimator robust to a wide class of omitted variable problems that would plague pooled OLS.

(d) [16 points] Now, in order to implement your strategy discussed above, go to Yahoo Finance website and collect the stock price data of any ten U.S. companies of your choice as well as the stock price for S&P500 (i.e., as your measure of overall market performance and its ticker symbol is ˆGSPC) for the above sample. Use the adjusted close price (column) as your price data for your $n = 10$ firms and compute the percentage change of these series as your performance measure. Make sure that you arrange your data as a panel data format by stacking these 10 companies' dataset. Then (i) run a cross section fixed effect panel data model of firm performance on market performance only; and (ii) explain (i.e., not estimate) how you would estimate a model or control for time fixed effects as well.

SOLUTION:

Data Collection and Construction. Daily adjusted close prices were downloaded from Yahoo Finance for ten U.S. companies spanning diverse sectors and the S&P 500 index (`ˆGSPC`) for the period January 1, 2021 through January 1, 2026. The ten firms are:

ID	Ticker	Company	Sector
1	AAPL	Apple Inc.	Technology
2	MSFT	Microsoft Corp.	Technology
3	JPM	JPMorgan Chase & Co.	Financials
4	JNJ	Johnson & Johnson	Healthcare
5	XOM	ExxonMobil Corp.	Energy
6	PG	Procter & Gamble Co.	Consumer Staples
7	DIS	Walt Disney Co.	Communication Services
8	HD	Home Depot Inc.	Consumer Discretionary
9	CAT	Caterpillar Inc.	Industrials
10	NEE	NextEra Energy Inc.	Utilities

Daily log returns were computed as $r_{it} = 100 \times [\ln(P_{it}) - \ln(P_{i,t-1})]$, where P_{it} is the adjusted close price. This log-difference formulation is the standard econometric measure of percentage change (Stock and Watson 2020, Ch. 14). The first observation for each firm (which yields a missing return due to differencing) was dropped. The data were then stacked into a balanced panel with $n = 10$ firms $\times$ $T = 1{,}254$ trading days $= 12{,}540$ observations, identified by `firm_id` and `stata_date`.

(i) Cross-Section Fixed Effects Model. The following firm fixed effects model was estimated in Stata:

```
use "data/stock_panel.dta", clear
xtset firm_id stata_date
xtreg log_ret_firm log_ret_market, ///
    fe vce(robust)
```

The estimated model is:

$$\boxed{\widehat{\Delta \ln(P_{it})} = \hat{\alpha}_i + \underset{(0.110)}{0.782}\, \Delta \ln(P_{mt})}$$

$n = 12{,}540$ observations, 10 firms, $T = 1{,}254$ trading days per firm.

Within $R^2 = 0.2770$, Overall $R^2 = 0.2768$.

The estimated market beta is $\hat{\beta} = 0.782$ (SE $= 0.110$, $t = 7.11$, $p < 0.001$), indicating that a one-percentage-point increase in the S&P 500 daily log return is associated with, on average, a 0.782 percentage-point increase in the firm's daily log return, holding firm-specific factors constant. The within R^2 of 0.277 means that market movements explain about 27.7% of the within-firm day-to-day variation in individual stock returns, which is substantial for daily frequency data.

The estimated firm fixed effects ($\hat{\alpha}_i$) represent each firm's average daily idiosyncratic return after removing the market component—akin to "Jensen's alpha" in the CAPM framework (Jensen, 1968). For instance, ExxonMobil (XOM) and Caterpillar (CAT) have the largest positive fixed effects ($\hat{\alpha} \approx 0.056\%$ per day), while Disney (DIS) has the most negative ($\hat{\alpha} \approx -0.077\%$ per day), reflecting differential firm-level performance over the sample period.

See Table 6 in the Appendix for the full regression output.

(ii) Controlling for Time Fixed Effects. To additionally control for time (trading-day) fixed effects, one would augment the model with a full set of day-

specific dummy variables:

$$\Delta \ln(P_{it}) = \alpha_i + \delta_t + \beta \cdot \Delta \ln(P_{mt}) + u_{it}$$

where δ_t represents the fixed effect for trading day t. In Stata, one could either generate day dummies manually or use `i.stata_date` factor notation:

```
reg log_ret_firm log_ret_market ///
    i.firm_id i.stata_date,     ///
    vce(robust)
```

However, there is an important identification problem. The market return $\Delta \ln(P_{mt})$ varies only across time, not across firms—it takes the same value for all n firms on any given trading day t. Including a full set of day fixed effects δ_t (one dummy per trading day) would absorb all variation that is common to all firms on each day, which is exactly the variation that $\Delta \ln(P_{mt})$ exploits for identification. In the extreme, $\Delta \ln(P_{mt})$ is a deterministic function of δ_t (since it is perfectly collinear with the day dummies), and the coefficient β would be unidentified.

In practice, if we want to control for aggregate daily shocks while still estimating β, we have several options. One approach is to include **time fixed effects but drop** $\Delta \ln(P_{mt})$, letting the day dummies absorb all common daily variation; this controls for market movements non-parametrically but does not produce an estimate of β. A second approach is to use **year or month fixed effects** rather than daily fixed effects, which absorb low-frequency common trends while leaving enough within-period variation in $\Delta \ln(P_{mt})$ for β to be identified. A third

approach is to include additional **time-varying aggregate controls** (such as changes in the VIX volatility index, interest rate changes, or oil price changes) alongside $\Delta \ln(P_{mt})$, rather than saturating the model with day dummies.

The simplest implementation for time FE that avoids the collinearity issue would be:

```
* Year FE (coarse; avoids collinearity)
gen year = year(dofC(stata_date))
xtreg log_ret_firm log_ret_market ///
    i.year, fe vce(robust)
```

Day FE and $\Delta \ln(P_{mt})$ are collinear; use coarser time FE (year/month) or drop $\Delta \ln(P_{mt})$.

■

Reflections: The collinearity between time fixed effects and aggregate regressors is a fundamental identification issue in panel data. When the regressor of interest varies only across time (like $\Delta \ln(P_{mt})$), entity fixed effects are fine (they remove cross-sectional variation that the regressor does not exploit), but time fixed effects remove the very source of variation that identifies the coefficient. This is a concrete illustration of the general principle that fixed effects "eat" the variation they span: including T time dummies leaves only cross-sectional variation within each period, and if the regressor is identical across entities within each period, there is no variation left.

The recommended approach depends on the research question. If the goal is to estimate market beta (β), use firm FE without day FE (or with coarse time

FE). If the goal is to isolate firm-specific shocks from all aggregate factors, use two-way FE (firm + day) but recognize that β is not separately identified from δ_t (Stock and Watson 2020, Ch. 10).

Ungraded Exercises

The following questions will not be graded; they are for practice and will be discussed at the recitation.

SW Exercise 10.1

This exercise refers to the drunk driving panel data regression summarized in Table 10.1.

(a) New Jersey has a population of 8.85 million people. Suppose New Jersey increased the tax on a case of beer by $2 (in 1988 dollars). Use the results in column (5) to predict the number of lives that would be saved over the next year. Construct a 99% confidence interval for your answer.

(b) The drinking age in New Jersey is 21. Suppose that New Jersey lowered its drinking age to 19. Use the results in column (5) to predict the change in the number of traffic fatalities in the next year. Construct a 95% confidence interval for your answer.

(c) Suppose real income per capita in New Jersey increases by 3% in the next year. Use the results in column (6) to predict the change in the number of traffic fatalities in the next year. Construct a 95% confidence interval for your answer.

(d) How should standard errors be clustered in the regressions in columns (2) through (7)?

(e) How should minimum drinking age be included in the regressions? Should it enter as a continuous variable or as a series of indicator variables? Be specific about the information you use to assess this question.

Solution:

This exercise uses the regressions in Table 10.1 (Stock and Watson 2020, p. 336), which reports panel regressions of the traffic fatality rate (deaths per 10,000 people) on beer taxes, drinking-age laws, and other controls for 48 U.S. states over 1982–1988, with state and year fixed effects and clustered standard errors.

(a) From column (5), the coefficient on *Beer tax* is $\hat{\beta} = -0.69$ with SE $= 0.35$ and 95% CI $[-1.38, 0.00]$. The dependent variable is the fatality rate per 10,000 people. The predicted change in the fatality rate from a \$2 tax increase is $\Delta \widehat{FatalityRate} = -0.69 \times 2 = -1.38$ deaths per 10,000. For New Jersey (population 8.85 million = 885 units of 10,000), the predicted number of lives saved is

$$\Delta \widehat{Deaths} = -1.38 \times 885 = -1221.3.$$

For the 99% CI, note that $\Delta \widehat{Deaths} = \hat{\beta} \times 2 \times 885 = 1770 \times \hat{\beta}$. The 99% CI for $\hat{\beta}$ is $\hat{\beta} \pm 2.576 \times \text{SE} = -0.69 \pm 2.576(0.35) = [-1.59, 0.21]$. Multiplying by 1770:

$$\boxed{\text{99\% CI for lives saved} = [1770 \times (-1.59),\ 1770 \times 0.21] = [-2814,\ 372].}$$

The point estimate predicts about 1,221 lives saved, but the confidence interval includes positive values (i.e., more deaths), reflecting the imprecision of the beer tax coefficient.

(b) From column (5), the drinking-age variables enter as indicators for ages 18, 19, and 20 (with age 21 as the omitted category). To predict the effect of lowering the drinking age from 21 to 19, we use the coefficient on "Drinking age 19," which is $\hat{\delta}_{19} = -0.08$ with SE $= 0.07$ and 95% CI $[-0.21, 0.06]$. The predicted change in the fatality rate is

$\Delta \widehat{FatalityRate} = -0.08$ per 10,000. For New Jersey: $\Delta \widehat{Deaths} = -0.08 \times 885 = -70.8$. The 95% CI is $[885 \times (-0.21),\ 885 \times 0.06]$:

$$\boxed{\text{95\% CI for change in deaths} = [-185.9,\ 53.1].}$$

The point estimate suggests about 71 fewer deaths (counterintuitively), but the coefficient is statistically insignificant and the confidence interval includes both negative and positive values. The negative sign likely reflects that variation in drinking-age laws is limited and the effect is imprecisely estimated once state and year fixed effects are included.

(c) From column (6), the coefficient on *Real income per capita (logarithm)* is $\hat{\beta}_{inc} = 1.79$ with SE $= 0.64$ and 95% CI $[1.79 \pm 1.96(0.64)] = [0.54, 3.04]$. Because the regressor is $\ln(income)$, a 3% increase in real income corresponds to $\Delta \ln(income) = \ln(1.03) \approx 0.0296$. The predicted change in the fatality rate is $\Delta \widehat{FatalityRate} = 1.79 \times 0.0296 = 0.0530$ per 10,000. For New Jersey:

$$\Delta \widehat{Deaths} = 0.0530 \times 885 \approx 46.9.$$

The 95% CI is $[0.54 \times 0.0296 \times 885,\ 3.04 \times 0.0296 \times 885]$:

$$\boxed{\text{95\% CI} = [14.2,\ 79.7] \text{ additional deaths.}}$$

Higher income is associated with more traffic deaths, plausibly because higher income leads to more driving. The effect is statistically significant.

(d) Standard errors should be **clustered at the state level**. Each state is observed

over multiple years, and the error terms for a given state are likely serially correlated (e.g., a state with unusually dangerous roads in 1984 will probably still have dangerous roads in 1985). Clustering at the state level allows for arbitrary within-state correlation over time and for heteroskedasticity across states. This is the approach recommended by Bertrand et al. (2004) and implemented in all columns (2)–(7) of Table 10.1.

(e) The minimum drinking age should enter as a **series of indicator (dummy) variables** rather than as a continuous variable. The reason is that Table 10.1 provides evidence to assess this question directly. In columns (4) and (5), the drinking age enters as three dummies (age 18, 19, 20), with age 21 as the omitted category. The coefficients are 0.03, −0.02, and 0.03 in column (4) and −0.01, −0.08, −0.10 in column (5)—they do not follow a monotonic linear pattern. In column (6), the drinking age enters as a single continuous variable with a coefficient of 0.00 (SE = 0.02), which is completely insignificant. The F-test for the joint significance of the three drinking-age dummies yields $F = 0.35$ ($p = 0.786$) in column (4) and $F = 1.41$ ($p = 0.253$) in column (5), both insignificant. The continuous specification masks the potentially nonlinear relationship; thus, indicator variables are preferred because they allow the effect of each age level to be estimated freely without imposing linearity. ■

Reflections: Exercise 10.1 reinforces several core panel data concepts. The key calculation skill is translating regression coefficients into real-world predictions: multiplying the coefficient by the change in the regressor and then scaling by the population. The log-level specification in part (c) requires computing $\Delta \ln(income) = \ln(1.03)$ rather than simply using 0.03. Part (d) reinforces that clustering at the entity level is essential in panel data to account for serial correlation (Arellano 1987; Bertrand et al. 2004). Part (e) illustrates the "continuous vs. dummies" modeling choice: when the

relationship between the regressor and the dependent variable may be nonlinear, indicator variables are more flexible, and the F-test for joint significance of the dummies tests whether the variable matters at all.

SW Exercise 10.5

Consider the model with a single regressor. This model also can be written as

$$\begin{aligned} Y_{it} = \beta_0 + \beta_1 X_{1,it} + \delta_2 B2_t + \cdots + \delta_T BT_t \\ + \gamma_2 D2_i + \cdots + \gamma_n Dn_i + u_{it}, \end{aligned}$$

where $B2_t = 1$ if $t = 2$ and 0 otherwise, $D2_i = 1$ if $i = 2$ and 0 otherwise, and so forth. How are the coefficients $(\beta_0, \delta_2, \ldots, \delta_T, \gamma_2, \ldots, \gamma_n)$ related to $(\alpha_1, \ldots, \alpha_n, \lambda_1, \ldots, \lambda_T)$?

Solution:

The entity and time fixed effects model can be written in two equivalent ways. The first uses entity-specific intercepts α_i and time-specific intercepts λ_t:

$$Y_{it} = \beta_1 X_{1,it} + \alpha_i + \lambda_t + u_{it}, \qquad i = 1, \ldots, n, \quad t = 1, \ldots, T.$$

The second uses a common intercept β_0 and deviations from the reference categories via dummy variables:

$$Y_{it} = \beta_0 + \beta_1 X_{1,it} + \delta_2 B2_t + \cdots + \delta_T BT_t$$

$$+ \gamma_2 D2_i + \cdots + \gamma_n Dn_i + u_{it},$$

where entity 1 and time period 1 are the omitted reference categories. To derive the mapping, evaluate the second equation for each combination:

For entity $i = 1$, time $t = 1$: all dummies are zero, so $E[Y_{1,1} \mid X] = \beta_0 + \beta_1 X_{1,1,1}$. From the first formulation, this equals $\alpha_1 + \lambda_1 + \beta_1 X_{1,1,1}$. Therefore:

$$\beta_0 = \alpha_1 + \lambda_1.$$

For entity $i = 1$, time $t = s$ (where $s \geq 2$): only $B_s = 1$, so $E[Y_{1,s} \mid X] = \beta_0 + \delta_s + \beta_1 X_{1,1,s} = (\alpha_1 + \lambda_1) + \delta_s + \beta_1 X_{1,1,s}$. From the first formulation, this equals $\alpha_1 + \lambda_s + \beta_1 X_{1,1,s}$. Equating:

$$\delta_s = \lambda_s - \lambda_1, \qquad s = 2, \ldots, T.$$

For entity $i = j$ (where $j \geq 2$), time $t = 1$: only $D_j = 1$, so $E[Y_{j,1} \mid X] = \beta_0 + \gamma_j + \beta_1 X_{1,j,1} = (\alpha_1 + \lambda_1) + \gamma_j + \beta_1 X_{1,j,1}$. From the first formulation, this equals $\alpha_j + \lambda_1 + \beta_1 X_{1,j,1}$. Equating:

$$\gamma_j = \alpha_j - \alpha_1, \qquad j = 2, \ldots, n.$$

Collecting these results:

$$\boxed{\begin{aligned} &\beta_0 = \alpha_1 + \lambda_1 \\ &\delta_s = \lambda_s - \lambda_1, \quad s = 2, \ldots, T \\ &\gamma_j = \alpha_j - \alpha_1, \quad j = 2, \ldots, n \end{aligned}}$$

The common intercept β_0 absorbs the fixed effects of the reference entity (entity 1) and the reference period (period 1). Each time dummy coefficient δ_s measures the *deviation* of period s from the reference period. Each entity dummy coefficient γ_j measures the *deviation* of entity j from the reference entity. The slope coefficient β_1 is identical in both formulations, as it captures the within-entity, within-period effect of X on Y regardless of the parameterization. ■

Reflections: This exercise clarifies an important but often confusing point about the two equivalent parameterizations of the TWFE model. In the "$\alpha_i + \lambda_t$" formulation, one entity and one time effect must be normalized (e.g., set to zero) to avoid perfect multicollinearity with the intercept. Stata's `i.entity i.time` notation automatically drops one category from each set. The dummy-variable formulation with β_0 makes the normalization explicit: entity 1 and period 1 serve as the reference categories, and all other effects are measured relative to them. The key insight is that only the *differences* $\alpha_j - \alpha_1$ and $\lambda_s - \lambda_1$ are identified—the absolute levels of the fixed effects are not separately identified from the intercept. This does not affect estimation of β_1, which is the parameter of interest (Stock and Watson 2020, Ch. 10; Wooldridge 2019, Ch. 14).

SW Empirical Exercise E10.2

Do citizens demand more democracy and political freedom as their incomes grow? That is, is democracy a normal good? On the text website, you will find the data file **Income_Democracy**, which contains a panel data set from 195 countries for the years 1960, 1965, ..., 2000. A detailed description is given in **Income_Democracy_Description**, available on the website. The data set contains an index of political freedom/democracy for each country in each year, together with data on each country's income and various demographic controls. (The income and demographic controls are lagged five years relative to the democracy index to allow time for democracy to adjust to changes in these variables.)

(a) Is the data set a balanced panel? Explain.

(b) The index of political freedom/democracy is labeled *Dem_ind*.

  (i) What are the minimum and maximum values of *Dem_ind* in the data set? What are the mean and standard deviation of *Dem_ind* in the data set? What are the 10th, 25th, 50th, 75th, and 90th percentiles of its distribution?

  (ii) What is the value of *Dem_ind* for the United States in 2000? Averaged over all years in the data set?

  (iii) What is the value of *Dem_ind* for Libya in 2000? Averaged over all years in the data set?

  (iv) List five countries with an average value of *Dem_ind* greater than 0.95; less than 0.10; and between 0.3 and 0.7.

(c) The logarithm of per capita income is labeled *Log_GDPPC*. Regress *Dem_ind* on *Log_GDPPC*. Use standard errors that are clustered by country.

(i) How large is the estimated coefficient on *Log_GDPPC*? Is the coefficient statistically significant?

(ii) If per capita income in a country increases by 20%, by how much is *Dem_ind* predicted to increase? What is a 95% confidence interval for the prediction? Is the predicted increase in *Dem_ind* large or small? (Explain what you mean by large or small.)

(iii) Why is it important to use clustered standard errors for the regression? Do the results change if you do not use clustered standard errors?

(d) (i) Suggest a variable that varies across countries but plausibly varies little—or not at all—over time and that could cause omitted variable bias in the regression in (c).

(ii) Estimate the regression in (c), allowing for country fixed effects. How do your answers to (c)(i) and (c)(ii) change?

(iii) Exclude the data for Azerbaijan, and rerun the regression. Do the results change? Why or why not?

(iv) Suggest a variable that varies over time but plausibly varies little—or not at all—across countries and that could cause omitted variable bias in the regression in (c).

(v) Estimate the regression in (c), allowing for time and country fixed effects. How do your answers to (c)(i) and (c)(ii) change?

(vi) There are additional demographic controls in the data set. Should these variables be included in the regression? If so, how do the results change when they are included?

(e) Based on your analysis, what conclusions do you draw about the effects of income on

democracy?

SOLUTION:

The dataset `income_democracy.dta` contains a panel of up to 195 countries observed every five years from 1960 to 2000 (9 time periods). The key variables are `Dem_ind` (a democracy index scaled 0 to 1) and `Log_GDPPC` (the log of real GDP per capita, lagged 5 years). The Stata commands and results below draw on the analysis in Recitation 9.

```
use "income_democracy.dta", clear
xtset code year, delta(5)
```

(a) A **balanced panel** requires every cross-sectional unit to be observed in every time period. Tabulating observations by year reveals that N ranges from 95 in 1960 to 195 in 2000, so the panel is **unbalanced**. New countries enter the dataset over time as former colonies gain independence and post-Soviet states appear after 1990.

(b)(i) Using `summarize dem_ind, detail`:

Min	0	Mean	0.481
p10	0	Std. Dev.	0.372
p25	0.167	p75	0.833
p50 (median)	0.500	p90	1
Max	1		

The distribution is roughly bimodal: many country-years cluster near 0 (autocracies) or 1 (democracies), with the median at 0.5.

(b)(ii–iii) The United States scores 1.0 in 2000 with an all-years average of 0.986 (slightly below 1 because early years capture pre–civil-rights-era restrictions). Libya scores 0.0 in 2000 with an all-years average of 0.109 under the Gaddafi regime.

(b)(iv)

Avg. > 0.95	Avg. < 0.10	Avg. $\in [0.3, 0.7]$
Australia	Afghanistan	Argentina
Canada	China	Brazil
Denmark	Cuba	Chile
Norway	Iraq	Mexico
United States	Saudi Arabia	Turkey

(c)(i) The pooled OLS regression $\texttt{Dem_ind}_{it} = \beta_0 + \beta_1 \texttt{Log_GDPPC}_{it} + u_{it}$ with clustered standard errors yields $\hat{\beta}_1 = 0.236$ (SE $= 0.012$, $t = 19.91$, $p < 0.001$). The coefficient is highly statistically significant: a one-unit increase in log GDP per capita (roughly a 170% increase in income) is associated with a 0.236-point increase in the democracy index.

(c)(ii) A 20% increase in income corresponds to $\Delta \ln(GDPPC) = \ln(1.20) \approx 0.1823$. The predicted change is $\Delta\widehat{\texttt{Dem_ind}} = 0.236 \times 0.1823 \approx 0.043$. The 95% CI for $\hat{\beta}_1$ is $[0.212, 0.259]$, so the 95% CI for $\Delta\texttt{Dem_ind}$ is $[0.212 \times 0.1823,\ 0.259 \times 0.1823] = [0.039, 0.047]$. On a 0-to-1 scale, an increase of ≈ 0.043 is **relatively small**: a substantial 20% income increase closes less than 5% of the gap between a fully autocratic and fully democratic country.

(c)(iii) Clustering is essential because panel observations within each country are serially correlated: a country that is more democratic than predicted in 1980 is likely

still so in 1985. Without clustering, standard errors are too small (SE = 0.0086 vs. 0.012 clustered), leading to inflated t-statistics and spuriously precise confidence intervals. The point estimate is unaffected by clustering—only the standard errors change.

(d)(i) A strong candidate is a country's **colonial/legal origin** (e.g., British common law vs. French civil law). This is essentially time-invariant, varies across countries, and is correlated with both income (former British colonies tend to be richer) and democracy (they inherited democratic institutions). Omitting it biases the pooled OLS upward.

(d)(ii) With country fixed effects (`xtreg dem_ind log_gdppc, fe cluster(code)`), the coefficient drops dramatically from 0.236 to $\hat{\beta}_1 = 0.084$ (SE = 0.032, $t = 2.66$, $p = 0.009$)—a 64% reduction. The effect of a 20% income increase falls from 0.043 to $0.084 \times 0.1823 \approx 0.015$. Most of the pooled OLS correlation was driven by between-country differences (rich countries are democracies), not within-country income growth causing democratization.

(d)(iii) Excluding Azerbaijan has **no effect** on the results ($\hat{\beta}_1 = 0.084$, n drops from 958 to 957). Azerbaijan has only one observation (year 2000), so the within estimator—which requires over-time variation—assigns it zero weight. Its demeaned values are identically zero. Any country observed in only one period is automatically irrelevant to the FE estimator.

(d)(iv) A natural candidate is a **global democratization wave** (the international political "climate" in each period). This varies over time (decolonization in the 1960s, the third wave of democratization in the 1980s–90s, the end of the Cold War) but is roughly constant across countries in any given year. It is correlated with both global

income growth and democracy, creating omitted variable bias. The fix is to include time fixed effects (year dummies).

(d)(v) With two-way FE (entity + time: `xtreg dem_ind log_gdppc i.year, fe cluster(code)`), the coefficient drops further to $\hat{\beta}_1 = 0.054$ (SE $= 0.042$, $t = 1.26$, $p = 0.209$). The 95% CI is $[-0.030, 0.137]$, which **includes zero**. The effect of a 20% income increase is now $0.054 \times 0.1823 \approx 0.010$—less than one percentage point and statistically insignificant. The year F-test yields $F = 5.73$ ($p < 0.001$), confirming that time effects are jointly significant.

(d)(vi) Adding demographic controls (age shares, education, log population) to the two-way FE model further attenuates the coefficient to $\hat{\beta}_1 = 0.025$ (SE $= 0.054$, $t = 0.47$, $p = 0.642$). The 95% CI is $[-0.082, 0.132]$. The F-test for joint significance of all demographic controls yields $F = 1.44$ ($p = 0.21$)—they are not jointly significant. The sample drops from 958 to 680 due to missing education data.

(e) The progressive results are summarized in the table below:

Specification	Controls	$\hat{\beta}_1$	Significant?
Pooled OLS	—	0.236	Yes ($p < 0.001$)
+ Time FE	—	0.235	Yes ($p < 0.001$)
+ Entity FE	—	0.084	Yes ($p = 0.009$)
+ Entity & Time FE	—	0.054	No ($p = 0.209$)
+ Entity & Time FE	Demographics	0.025	No ($p = 0.642$)

The strong cross-country correlation between income and democracy ($\hat{\beta}_1 = 0.236$ in pooled OLS) is largely spurious, driven by time-invariant confounders such as colonial history, geography, and culture. Once country fixed effects are included, the coefficient

drops by 64%, and once time fixed effects are added, it becomes statistically insignificant. In the most demanding specifications (two-way FE $\pm$ controls), the effect of income on democracy is small, insignificant, and its confidence interval includes zero. A 20% increase in per capita income—substantial economic growth—is predicted to raise the democracy index by less than one percentage point, and we cannot reject a true effect of zero.

> Democracy does not appear to be a "normal good" in the causal sense. The income–democracy correlation reflects common historical causes, not a direct causal channel from economic development to political liberalization.

This conclusion is consistent with Stock and Watson (2020), Ch. 10, and the influential findings of Acemoglu, Johnson, Robinson, and Yared (2008, *American Economic Review*), who argue that the income–democracy relationship reflects joint determination by historical institutional development rather than a causal "modernization" effect. ■

Reflections: Empirical Exercise E10.2 is the single most instructive exercise in the chapter because it demonstrates the entire logic of panel data econometrics in a single application. The key lesson is that the cross-sectional (between-country) correlation between income and democracy is large and statistically significant, but it is almost entirely attributable to omitted time-invariant confounders. As we progressively add entity FE, time FE, and controls, the coefficient shrinks from 0.236 to 0.025 and loses significance. This pattern is strikingly parallel to the concealed-carry analysis in Problems 1–2 of this problem set, where the shall-issue coefficient shrank from -0.443 to -0.028 as fixed effects were added. In both cases, the moral is the same: pooled

OLS can be severely misleading when unobserved heterogeneity is present, and fixed effects estimation is essential for credible causal inference in panel data.

The Azerbaijan example is a subtle but important pedagogical point: entities observed in only one time period contribute zero information to the within estimator, because there is no within-entity variation to exploit. This is a fundamental feature of the FE approach and a key difference from pooled OLS, which uses all observations regardless of how many periods each entity contributes.

Appendix: Tables and Figures

Table 4: Problem 1 Regression (i): ln(vio) on shall

	(1) ln(Violent Crime Rate)
shall	-0.443^{***}
	(0.0475)
Constant	6.135^{***}
	(0.0193)
Observations	$1,173$
R^2	0.0866
Adjusted R^2	0.0859

Standard errors in parentheses

Heteroskedasticity-robust standard errors in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Table 5: Problem 1 Regression (ii): ln(vio) on shall + controls

	(1) ln(Violent Crime Rate)
shall	-0.368^{***}
	(0.0348)
incarc_rate	0.00161^{***}
	(0.000181)
density	0.0267^{*}
	(0.0143)
avginc	0.00121
	(0.00728)
pop	0.0427^{***}
	(0.00315)
pb1064	0.0809^{***}
	(0.0200)
pw1064	0.0312^{***}
	(0.00973)
pm1029	0.00887
	(0.0121)
Constant	2.982^{***}
	(0.609)
Observations	$1,173$
R^2	0.5643
Adjusted R^2	0.5613

Standard errors in parentheses

Heteroskedasticity-robust standard errors in parentheses.

Controls: incarc_rate, density, avginc, pop, pb1064, pw1064, pm1029.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Table 6: Firm-Level Log Returns on S&P 500 Log Returns: Fixed Effects Panel Regression

	(1) Log return of firm stock (x100)
Log return of S&P 500 (x100)	0.782***
	(0.110)
Constant	0.00905
	(0.00540)
Observations	12,540
Within R^2	0.2770
Between R^2	
Overall R^2	0.2768
Number of firms	10

Standard errors in parentheses

Dependent variable: daily log return of firm stock (x100).

Log returns: $100 \times [\ln(P_it) - \ln(P_i, t-1)]$.

Firm (cross-section) fixed effects via within transformation.

Heteroskedasticity-robust standard errors in parentheses.

Data: Yahoo Finance, Jan 2021 – Dec 2025, 10 U.S. firms.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Replication Files

All replication materials for `homework_05` are available on GitHub:

- Course repository: https://github.com/tylersotomayor/econun3412-spring2026
- Homework 05 folder: https://github.com/tylersotomayor/econun3412-spring2026/tree/main/homework_05

References

Arellano, M. (1987). Computing robust standard errors for within-groups estimators. *Oxford Bulletin of Economics and Statistics*, 49(4):431–434.

Bertrand, M., Duflo, E., and Mullainathan, S. (2004). How much should we trust differences-in-differences estimates? *Quarterly Journal of Economics*, 119(1):249–275.

Dimson, E. (1979). Risk measurement when shares are subject to infrequent trading. *Journal of Financial Economics*, 7(2):197–226.

Donohue, III, J. J. and Ayres, I. (2003). Shooting down the "More guns, less crime" hypothesis. *Stanford Law Review*, 55(4):1193–1312.

Fama, E. F. and French, K. R. (1993). Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics*, 33(1):3–56.

Fama, E. F. and MacBeth, J. D. (1973). Risk, return, and equilibrium: Empirical tests. *Journal of Political Economy*, 81(3):607–636.

Jegadeesh, N. and Titman, S. (1993). Returns to buying winners and selling losers: Implications for stock market efficiency. *Journal of Finance*, 48(1):65–91.

Jensen, M. C. (1968). The performance of mutual funds in the period 1945–1964. *Journal of Finance*, 23(2):389–416.

Lintner, J. (1965). The valuation of risk assets and the selection of risky investments in stock portfolios and capital budgets. *Review of Economics and Statistics*, 47(1):13–37.

Lott, Jr., J. R. and Mustard, D. B. (1997). Crime, deterrence, and right-to-carry concealed handguns. *Journal of Legal Studies*, 26(1):1–68.

National Research Council (2005). Firearms and violence: A critical review. Technical report, The National Academies Press, Washington, DC.

Sharpe, W. F. (1964). Capital asset prices: A theory of market equilibrium under conditions of risk. *Journal of Finance*, 19(3):425–442.

Stock, J. H. and Watson, M. W. (2020). *Introduction to Econometrics*. Pearson, New York, 4th edition.

Wooldridge, J. M. (2019). *Introductory Econometrics: A Modern Approach*. Cengage Learning, Boston, 7th edition.